

Rahul Yadav

+1(716)238-0835 | rahulyad@buffalo.edu | [linkedin.com/in/rahul-yadav-704](https://www.linkedin.com/in/rahul-yadav-704) | [Github](#) | [Website](#)

EDUCATION

University at Buffalo, Buffalo <i>Master of Science in Computer Science and Engineering</i> Coursework: NLP, Computational Linguistics, Pattern Recognition, Machine Learning, Algorithms Design	Aug. 2022 – Dec. 2023 New York, US GPA: 3.82 / 4
Indian Institute of Technology (Banaras Hindu University), Varanasi <i>Bachelor of Technology in Electronics Engineering</i>	Jul. 2014 – May 2018 Uttar Pradesh, India

EXPERIENCE

Hilabs <i>Senior Data Scientist Applied Data Science Team</i>	Feb. 2024 – May 2024 Washington, D.C.
---	--

GenAI Driven Contract Analyzer: Streamlining Claim Processing <ul style="list-style-type: none">Led development of scalable end-to-end system for processing provider contract documents, emphasizing entity extraction to facilitate pricing configuration in claim processing applicationsEngineered custom document processing pipeline with LayoutLM extracting entities with images, tables, and text elementsDeveloped Langchain service for indexing extracted elements and managed metadata within OpenSearch to optimize retrievalImplemented a robust infrastructure leveraging self-hosted fine-tuned Mistral-AI LLM and asynchronous processing via SQS to efficiently handle requests, while prioritizing the utmost security and privacy of sensitive legal contractsSuccessful POCs followed by integration for 3 clients to validate solution effectiveness and potential for widespread adoption	
--	--

Apexanalytix <i>Data Scientist ML Engineer Intern Analytics Team</i>	May. 2023 – Feb. 2024 Remote, New York
--	---

Generative Knowledge Specific Chatbot <ul style="list-style-type: none">Developed advanced Retrieval Augmented Generation(RAG) Chatbot with LLM for intelligent knowledge access across teamsEngineered efficient retrieval pipeline with Parent Child document indexing using LangChain and Chroma VectorDBDesigned agile RAG pipeline, integrating MMR scoring for chunk retrieval and Azure OpenAI LLM for response generationImplemented user feedback collection to monitor chatbot performance and gather data for iterative refinement and fine-tuningAchieved 89% approval in human evaluations and integrated technology into 12 internal and 18 external client applications	
---	--

Oracle <i>Senior Application Engineer(ML) Oracle Financial Machine Learning Team</i>	Sept. 2020 – Aug. 2022 Bengaluru, India
--	--

Preemptive Anomaly Prediction in Corporate Billing <ul style="list-style-type: none">Implemented in-memory multivariate anomaly prediction system for corporate billing, addressing monthly billing challengesLeveraged Oracle in-database ML for Semi-Supervised classification with both local and global model explainabilityOptimized service with indexing and parallelism, processing 1.2M bills and 5M segments in 20 mins with 92% precisionIntegrated services with Oracle Revenue Management and Billing (ORMB) product, USPTO patent granted [US17/710745]	
--	--

Wipro Limited <i>Project Engineer(AI) AI Research Team</i>	June 2018 – Sept 2020 Bengaluru, India
--	---

Chatbot Services for Employee Helpline Portal's Ticketing System <ul style="list-style-type: none">Developed Employee Helpline Portal Chatbot with effective retrieval of historical ticket resolutions for enhanced user supportImplemented query intent classification and BERT-powered semantic search to deliver precise responsesSuccessfully integrated the chatbot into the portal and achieved the target human agent intervention reduction of 70%Achieved 79% accuracy score in evaluation and decreased wait times from 18 to 4 minutes, improving overall user experience	
---	--

PATENTS and PUBLICATIONS

Paper Published: Virtual Conversation with Real-Time Prediction of Body Moments/Gestures <i>Gopichand Agnihotram, Rajesh Kumar, Pandurang Naik, Rahul Yadav</i>	ICMLIP 2019 springerLink
Patent Granted: Method And System For Multimodal Analysis Based Emotion Recognition	US16/795840 [Link]
Patent Granted: Technology System For Assisting Financial Institutions In Debt Collection	US17/659017 [Link]

SKILLS

Languages : Python, Java, C++, C, PLSQL, SQL, MongoDB
Technologies : Retrieval Augmented Generation(RAG), Large Language Models(LLMs), Natural Language Processing(NLP), Machine Learning, Deep Learning, Web Services, Data Structures, Algorithms, Prompt Engineering, Indexing, Quantization
Cloud : Amazon Web Services (AWS), Azure, Google Cloud Platform (GCP), Oracle, Atlas, OpenAI
Frameworks and Libraries : Pytorch, Tensorflow, MLFlow, Kafka, SentenceTransformers(embeddings, re-rankers), Databricks, Snowflake, PySpark, Langchain, Llamaindex, VectorDBs(ChromaDB, Faiss, elasticsearch), MLFlow, Git, Jenkins, Big Data, Hadoop, Flask, Docker, Kubernetes, Pandas, Streamlit, Flask-RESTful, FastAPI, YOLO, Django, XGBoost, GAN, ActiveMQ, Springboot